

Face Recognition Vendor Test Ongoing

General Evaluation Specifications

VERSION 1.2

Patrick Grother
Mei Ngan
Kayee Hanaoka
Information Access Division
Information Technology Laboratory

Contact via frvt@nist.gov

February 14, 2022

NIST
National Institute of
Standards and Technology
U.S. Department of Commerce

Revision History

Date	Version	Description
April 1, 2019	1.0	Initial document
September 9, 2020	1.1	Update operating system to CentOS 8.2 and compiler to g++ 8.3.1 Adjust the legal similarity score range
February 14, 2022	1.2	Update operating system to Ubuntu 20.04.3 and compiler to g++ 9.3.0

Table of Contents

6		
7	1. Audience	3
8	2. Rules for Participation.....	3
9	2.1. Participation Agreement	3
10	2.2. Validation	3
11	2.3. Number and Schedule of Submissions	3
12	3. Reporting.....	3
13	3.1. Version Control.....	3
14	4. Hardware specification	4
15	5. Operating system, compilation, and linking environment	4
16	6. Software and Documentation.....	5
17	6.1. Library and Platform Requirements	5
18	6.2. Configuration and developer-defined data.....	5
19	6.3. A Note on Training.....	5
20	6.4. Submission folder hierarchy.....	5
21	6.5. Installation and Usage	6
22	6.6. Documentation.....	6
23	6.7. Modes of operation.....	6
24	7. Runtime behavior	6
25	7.1. Interactive behavior, stdout, logging	6
26	7.2. Exception Handling.....	6
27	7.3. External communication	6
28	7.4. Stateless behavior	6
29	7.5. Single-thread Requirement/Parallelization.....	6
30	8. Data structures supporting the API	7
31	8.1. Face Images	7
32	8.2. Data structure for eye coordinates	7
33	8.3. Template Role	8
34	8.4. Data type for similarity scores.....	8
35	8.5. Data structure for return value of API function calls	8
36	8.6. Data structure for quality elements	9

37

List of Tables

39	Table 3 – Structure for a single image	7
40	Table 4 – Labels describing categories of Images.....	7
41	Table 5 – Structure for a set of images from a single person.....	7
42	Table 6 – Structure for a pair of eye coordinates	8
43	Table 7 – Labels describing template role	8
44	Table 8 – Enumeration of return codes	8
45	Table 7 – ReturnStatus structure	9
46	Table 8 – Enumeration of quality elements.....	9

47

48

49 1. Audience

50 Participation in FRVT is open to any organization worldwide. There is no charge for participation. The target audience is
 51 researchers and developers of FR algorithms. While NIST intends to evaluate stable technologies that could be readily
 52 made operational, the test is also open to experimental, prototype and other technologies. All algorithms **must** be
 53 submitted as implementations of the API defined in the specific test's API document.

54 2. Rules for Participation

55 2.1. Participation Agreement

56 A participant must properly follow, complete, and submit the FRVT Participation Agreement. This must be done once,
 57 either prior or in conjunction with the very first algorithm submission. It is not necessary to do this for each submitted
 58 implementation thereafter UNLESS there are major organizational changes to the submitting entity.

59
 60 **NOTE** If an organization updates their cryptographic signing key, they must send a new completed participation
 61 agreement submission for this evaluation, with the fingerprint of their public key.

62 2.2. Validation

63 Prior to submission, all participants must run their software through the provided corresponding validation package for
 64 the test they wish to enter. The validation package will be made available at <https://github.com/usnistgov/frvt>. The
 65 purpose of validation is to ensure consistent algorithm output between the participant's execution and NIST's execution.

66 2.3. Number and Schedule of Submissions

67 Participants may send one submission as often as every four calendar months from the last submission for evaluation.
 68 NIST will evaluate implementations on a first-come-first-served basis, and quickly publish results.

69 3. Reporting

70 Unless otherwise specified for a specific test, for all algorithms that complete the evaluations, NIST will post performance
 71 results on the NIST FRVT website. NIST will maintain an email list to inform interested parties of updates to the website.
 72 Artifacts will include a leaderboard highlighting the top performing submissions in various areas (e.g., accuracy, speed
 73 etc.) and individual implementation-specific report cards. NIST will maintain reporting on the two most recent algorithm
 74 submissions from any organization. In the event an algorithm is no longer operable (e.g., license expiration, etc.), that
 75 algorithm will be retired from the evaluation. Prior submission results will be archived but remain accessible via a public
 76 link.

77
 78 **Important:** This is an open test in which NIST will identify the algorithm and the developing organization. Algorithm
 79 results will be attributed to the developer. Results will be machine generated (i.e. scripted) and will include timing,
 80 accuracy and other performance results. These will be posted alongside results from other implementations. Results will
 81 be expanded and modified as additional implementations are tested, and as analyses are implemented. Results may be
 82 regenerated on-the-fly, usually whenever additional implementations complete testing, or when new analysis is added.
 83
 84 NIST may additionally report results in workshops, conferences, conference papers and presentations, journal articles and
 85 technical reports.

86 3.1. Version Control

87 Developers must submit a version.txt file in the doc/ folder that accompanies their algorithm – see Section 6.4. The string
 88 in this file should allow the developer to associate results that appear in NIST reports with the submitted algorithm. This
 89 is intended to allow end-users to obtain productized versions of the prototypes submitted to NIST. NIST will publish the
 90 contents of version.txt. NIST has previously published MD5 hashes of the core libraries for this purpose.
 91

92 4. Hardware specification

93 NIST intends to support high performance by specifying the runtime hardware beforehand. There are several types of
 94 computer blades that may be used in the testing. Each machine has at least 128 GB of memory. We anticipate that 16
 95 processes can be run without time slicing, though NIST will handle all multiprocessing work via `fork()`¹. Participant-
 96 initiated multiprocessing is not permitted.

97 All implementations shall use 64 bit addressing.

98 NIST intends to support highly optimized algorithms by specifying the runtime hardware. There are several types of
 99 computers that may be used in the testing. The following list gives some details about possible compute architectures:

- 100 • Dual Intel® Xeon® E5-2630 v4 CPU @ 2.20GHz (10 cores each)²
- 101 • Dual Intel® Xeon® E5-2680 v4 CPU @ 2.4GHz (14 cores each)²
- 102 • Dual Intel® Xeon® Gold 6140 CPU @ 2.30GHz (18 cores each)³

103

104 All timing tests will be measured on Xeon R CPU E5-2630 v4 @ 2.20GHz processors. FRVT tests will not support the use of
 105 Graphics Processing Units (GPUs).

106 5. Operating system, compilation, and linking environment

107 The operating system that the submitted implementations shall run on will be released as a downloadable file accessible
 108 from <https://nigos.nist.gov/evaluations/ubuntu-20.04.3-live-server-amd64.iso> which is the 64-bit version of Ubuntu
 109 20.04.3 LTS (Focal Fossa) running Linux kernel 5.4.0-91-generic.

110 For this test, Windows machines will not be used. Windows-compiled libraries are not permitted. All software must run
 111 under Ubuntu 20.04.3.

112 NIST will link the provided library file(s) to our C++ language test drivers. Participants are required to provide their library
 113 in a format that is dynamically-linkable using the C++17 compiler g++ version 9.3.0.

114 A typical link line might be

115 `g++ -I. -Wall -m64 -o frvt11 frvt11.cpp -L. -lfrvt_11_acme_007`

116 The Standard C++ library should be used for development. Header files containing API function prototypes will be
 117 provided separately for each FRVT track and documented in each of the corresponding API documents.

118 The header files will be made available to implementers at <https://github.com/usnistgov/frvt>. All algorithm submissions
 119 will be compiled against the officially published header files – developers should not alter the header files when compiling
 120 and building their libraries.

121 All compilation and testing will be performed on x86_64 platforms. Thus, participants are strongly advised to verify
 122 library-level compatibility with g++ (on an equivalent platform) prior to submitting their software to NIST to avoid linkage
 123 problems later on (e.g. symbol name and calling convention mismatches, incorrect binary file formats, etc.).

124

¹ <http://man7.org/linux/man-pages/man2/fork.2.html>

² `cat /proc/cpuinfo` returns `fpu vme de pse tsc msr pae mce cx8 apic sep mtrr pge mca cmov pat pse36 clflush dts acpi mmx fxsr sse sse2 ss ht tm pbe syscall nx pdpe1gb rdtscp lm constant_tsc arch_perfmon pebs bts rep_good nopl xtopology nonstop_tsc aperfmperf eagerfpu pni pclmulqdq dtes64 monitor ds_cpl vmx smx est tm2 ssse3 fma cx16 xtpr pdcm pcid dca sse4_1 sse4_2 x2apic movbe popcnt tsc_deadline_timer aes xsave avx f16c rdrand lahf_lm abm 3dnowprefetch ida arat epb pln pts dtherm tpr_shadow vnmi flexpriority ept vpid fsgsbase tsc_adjust bmi1 hle avx2 smep bmi2 erms invpcid rtm cqm rdseed adx smap xsaveopt cqm_llc cqm_occup_llc`

³ `cat /proc/cpuinfo` returns `fpu vme de pse tsc msr pae mce cx8 apic sep mtrr pge mca cmov pat pse36 clflush dts acpi mmx fxsr sse sse2 ss ht tm pbe syscall nx pdpe1gb rdtscp lm constant_tsc arch_perfmon pebs bts rep_good nopl xtopology nonstop_tsc aperfmperf eagerfpu pni pclmulqdq dtes64 monitor ds_cpl vmx smx est tm2 ssse3 fma cx16 xtpr pdcm pcid dca sse4_1 sse4_2 x2apic movbe popcnt tsc_deadline_timer aes xsave avx f16c rdrand lahf_lm abm 3dnowprefetch ida arat epb pln pts dtherm tpr_shadow vnmi flexpriority ept vpid fsgsbase tsc_adjust bmi1 hle avx2 smep bmi2 erms invpcid rtm cqm mpx avx512f rdseed adx smap clflushopt avx512cd xsaveopt xsavec xgetbv1 xsaves cqm_llc cqm_occup_llc`

125 6. Software and Documentation

126 6.1. Library and Platform Requirements

127 Participants shall provide NIST with binary code only (i.e. no source code). The implementation should be submitted in
128 the form of a dynamically-linked library file.

129
130 The core library shall be named according to “Implementation Library Filename” documented in each API document. The
131 library name will generally follow the convention: *libfrvt_<track>_<provider>_<sequence>.so*. Additional supplemental
132 libraries may be submitted that support this “core” library file (i.e. the “core” library file may have dependencies
133 implemented in these other libraries). Supplemental libraries may have any name, but the “core” library must be
134 dependent on supplemental libraries in order to be linked correctly. The **only** library that will be explicitly linked to the
135 test driver is the “core” library.

136
137 Intel Integrated Performance Primitives (IPP) ® libraries are permitted if they are delivered as a part of the developer-
138 supplied library package. It is the provider’s responsibility to establish proper licensing of all libraries. The use of IPP
139 libraries shall not prevent running on CPUs that do not support IPP. Please take note that some IPP functions are
140 multithreaded and threaded implementations are prohibited.

141
142 Developers may obviously use common deep learning frameworks (e.g. Caffe, TensorFlow, etc.) and should submit those
143 dependencies as supplemental libraries. NIST has successfully received and run implementations leveraging such deep
144 learning frameworks in other evaluations with no issues.

145
146 Do not include any standard libraries (e.g., *libc.so*, *libgcc.so*, etc.) that come with the operating system and/or compilation
147 environment in your submission. The NIST test harness will handle all image I/O, so do not include JPEG or PNG libraries
148 (i.e., *libjpeg.so*, *libpng.so*) in your submission. If you need to include those libraries for other reasons, please contact NIST
149 prior to your submission. NIST will report the size of the supplied libraries.

150
151 **Important:** Public results will be attributed with the provider name and the 3-digit sequence number in the submitted
152 library name.

153 6.2. Configuration and developer-defined data

154 The implementation under test may be supplied with configuration files and supporting data files. NIST will report the
155 size of the supplied configuration files.

156 6.3. A Note on Training

157 NIST and the FRVT program do not train face recognition algorithms. We do not provide training data to software, and
158 software is prohibited from adapting to any data we pass to the algorithms. Training of face recognition algorithms is not
159 a turn-key operation; instead it is typically an extended process involving researchers curating suitable training sets,
160 establishing architectures and hyperparameters, and running trials over days or weeks, and then evaluating the
161 output. The result of such a process, which is often iterative, is usually a “trained model” i.e. static data and parameters
162 that can be saved and provided to NIST as an integral part of the black-box recognition engine. NIST does not support
163 training, because our tests seek to mimic operational reality and, there, algorithms are almost always shipped and used
164 “as is” without any training or adaptation to customer data. The representation of the face, as described by the “model”,
165 is fixed until the software is upgraded.

166 6.4. Submission folder hierarchy

167 Participant submissions shall contain the following folders at the top level

- 168 • *lib/* - contains all participant-supplied software libraries
- 169 • *config/* - contains all configuration and developer-defined data, e.g., trained models
- 170 • *doc/* - contains *version.txt*, which documents versioning information for the submitted software and any other
- 171 participant-provided documentation regarding the submission
- 172 • *validation/* - contains validation output

173 **6.5. Installation and Usage**

174 The implementation shall be installable using simple file copy methods. It shall not require the use of a separate
 175 installation program and shall be executable on any number of machines without requiring additional machine-specific
 176 license control procedures or activation. The implementation shall not use nor enforce any usage controls or limits based
 177 on licenses, number of executions, presence of temporary files, etc. The implementation shall remain operable for at
 178 least six months from the submission date.

179 **6.6. Documentation**

180 Participants shall provide documentation of additional functionality or behavior beyond that specified here.

181 **6.7. Modes of operation**

182 Implementations shall not require NIST to switch “modes” of operation or algorithm parameters. For example, the use of
 183 two different feature extractors must either operate automatically or be split across two separate library submissions.

184 **7. Runtime behavior**

185 **7.1. Interactive behavior, stdout, logging**

186 The implementation will be tested in non-interactive “batch” mode (i.e. without terminal support). Thus, the submitted
 187 library shall:

- 188 – Not use any interactive functions such as graphical user interface (GUI) calls, or any other calls which require
 189 terminal interaction e.g. reads from “standard input”.
- 190 – Run quietly, i.e. it should not write messages to "standard error" and shall not write to “standard output”.
- 191 – Only if requested by NIST for debugging, include a logging facility in which debugging messages are written to a
 192 log file whose name includes the provider and library identifiers and the process PID.

193 **7.2. Exception Handling**

194 The application should include error/exception handling so that in the case of a fatal error, the return code is still
 195 provided to the calling application.

196 **7.3. External communication**

197 Processes running on NIST hosts shall not side-effect the runtime environment in any manner, except for memory
 198 allocation and release. Implementations shall not write any data to external resource (e.g. server, file, connection, or
 199 other process), nor read from such, nor otherwise manipulate it. If detected, NIST will take appropriate steps, including
 200 but not limited to, cessation of evaluation of all implementations from the supplier, notification to the provider, and
 201 documentation of the activity in published reports.

202 **7.4. Stateless behavior**

203 All components in this test shall be stateless, except as noted. This applies to face detection, feature extraction and
 204 matching. Thus, all functions should give identical output, for a given input, independent of the runtime history. NIST
 205 will institute appropriate tests to detect stateful behavior. If detected, NIST will take appropriate steps, including but not
 206 limited to, cessation of evaluation of all implementations from the supplier, notification to the provider, and
 207 documentation of the activity in published reports.

208 **7.5. Single-thread Requirement/Parallelization**

209 Implementations must run in single-threaded mode, because NIST will parallelize the test by dividing the workload across
 210 many cores and many machines. Implementations must ensure that there are no issues with their software being
 211 parallelized via the `fork()` function. Developers should take caution with checking threading when using third-party
 212 frameworks (e.g., TensorFlow, MXNet, etc.).

8. Data structures supporting the API

The following section documents the common data structures used to support the C++ API functions for the various FRVT tasks. The actual C++ API function prototypes themselves are documented separately for each test and are available on the website for each track.

8.1. Face Images

An individual can be represented by $K \geq 1$ two-dimensional facial images.

Table 1 – Structure for a single image

C++ code fragment	Remarks
typedef struct Image	
{	
uint16_t image_width;	Number of pixels horizontally
uint16_t image_height;	Number of pixels vertically
uint16_t image_depth;	Number of bits per pixel. Legal values are 8 and 24.
std::shared_ptr<uint8_t> data;	Managed pointer to raster scanned data. Either RGB color or intensity. If image_depth == 24 this points to 3WH bytes RGBRGBRGB... If image_depth == 8 this points to WH bytes I I I I I I I I
Label description;	Single description of the image. The allowed values for this field are specified in the enumeration in Table 2.
} Image;	

An **Image** will be accompanied by one of the labels given below. Face recognition implementations should tolerate **Images** of any category.

Table 2 – Labels describing categories of Images

Label as C++ enumeration	Meaning
enum class Label {	
Unknown=0,	Either the label is unknown or unassigned.
Iso,	Frontal, intended to be in conformity to ISO/IEC 19794-5:2005.
Mugshot,	From law enforcement booking processes. Nominally frontal.
Photojournalism,	The image might appear in a news source or magazine. The images are typically taken by professional photographer and are well exposed and focused but exhibit pose and illumination variations.
Exploitation,	The image is taken from a child exploitation database. This imagery has highly unconstrained pose and illumination, expression and resolution.
Wild	Unconstrained image, taken by an amateur photographer, exhibiting wide variations in pose, illumination, and resolution.
};	

Table 3 – Structure for a set of images from a single person

C++ code fragment	Remarks
using Multiface = std::vector<Image>;	Vector of Image objects

8.2. Data structure for eye coordinates

The eye coordinates shall follow the placement semantics of the ISO/IEC 19794-5:2005 standard - the geometric midpoints of the endocanthion and exocanthion (see clause 5.6.4 of the ISO standard).

Sense: The label "left" refers to subject's left eye (and similarly for the right eye), such that $x_{right} < x_{left}$.

Table 4 – Structure for a pair of eye coordinates

C++ code fragment	Remarks
typedef struct EyePair	
{	
bool isLeftAssigned;	If the subject's left eye coordinates have been computed and assigned successfully, this value should be set to true, otherwise false.
bool isRightAssigned;	If the subject's right eye coordinates have been computed and assigned successfully, this value should be set to true, otherwise false.
uint16_t xleft;	X and Y coordinate of the center of the subject's left eye. If the eye coordinate is out of range (e.g. $x < 0$ or $x \geq \text{width}$), <code>isLeftAssigned</code> should be set to false, and the eye coordinates will be ignored.
uint16_t yleft;	
uint16_t xright;	X and Y coordinate of the center of the subject's right eye. If the eye coordinate is out of range (e.g. $x < 0$ or $x \geq \text{width}$), <code>isRightAssigned</code> should be set to false, and the eye coordinates will be ignored.
uint16_t yright;	
} EyePair;	

8.3. Template Role

Labels describing the type/role of the template to be generated will be provided as input to template generation.

Table 5 – Labels describing template role

Label as C++ enumeration	Meaning
enum class TemplateRole {	
Enrollment_11,	Enrollment template for 1:1 matching
Verification_11,	Verification template for 1:1 matching
Enrollment_1N,	Enrollment template for 1:N identification
Search_1N	Search template for 1:N identification
};	

8.4. Data type for similarity scores

Identification and verification functions shall return a measure of the similarity between the face data contained in the two templates. The datatype shall be an eight byte double precision real.

The similarity score values should be reported on the range that is used in the developer's software products. Larger values indicate more likelihood that the two samples are from the same person. However, we require scores to be non-negative. Developers often use [0,1], for example. Our test reports include various plots with threshold values e.g. FMR(T), to allow end-users to set thresholds in operations. These plots may become difficult to interpret if scores span many orders of magnitude.

Providers are cautioned that algorithms that natively produce few unique values will be disadvantaged by the inability to set a threshold precisely, as might be required to attain a false match rate of exactly 0.0001, for example.

8.5. Data structure for return value of API function calls

Table 6 – Enumeration of return codes

Return code as C++ enumeration	Meaning
enum class ReturnCode {	
Success=0,	Success
UnknownError,	Catch-all error
ConfigError,	Error reading configuration files
RefuseInput,	Elective refusal to process the input, e.g. because cannot handle greyscale
ExtractError,	Involuntary failure to process the image, e.g. after catching exception
ParseError,	Cannot parse the input data
TemplateCreationError,	Elective refusal to produce a template (e.g. insufficient pixels between the eyes)

VerifTemplateError,	For matching, either or both of the input templates were result of failed feature extraction
FaceDetectionError,	Unable to detect a face in the image
NumDataError,	The implementation cannot support the number of images
TemplateFormatError,	Template file is in an incorrect format or defective
EnrollDirError,	An operation on the enrollment directory failed (e.g., permission, space)
InputLocationError,	Cannot locate the input data – the input files or names seem incorrect
MemoryError,	Memory allocation failed (e.g., out of memory)
MatchError,	Error occurred during the 1:1 match operation
QualityAssessmentError,	Failure to generate a quality score on the input image
NotImplemented,	Function is not implemented
VendorError	Vendor-defined failure. Vendor errors shall return this error code and document the specific failure in the ReturnStatus.info string.
};	

Table 7 – ReturnStatus structure

C++ code fragment	Meaning
struct ReturnStatus {	
ReturnCode code;	Return Code
std::string info;	Optional information string
// constructors	
};	